

KnowLens — Workforce & Project Intelligence

Use Cases & Business Model

A governed AI intelligence layer that lets anyone in the org ask plain-language questions about work, projects, and people — answers scoped to exactly what they're allowed to see, with every query audited.

1. The system at a glance

Layer	Name	Role
Data store (System 1)	KnowCore	Source of truth — employees, projects, tasks, work logs, notes, progress summaries, workload signals
Intelligence layer	KnowLens	Governed MCP server — reads KnowCore with row-level scoping + aggregation; exposes tools/resources/prompts
Governance shell	Gateway (customized open-source MCP gateway)	Auth, tool-level RBAC, audit, discovery
Host	Claude	Where users actually ask questions

Positioning: a project-intelligence copilot with governed workforce analytics — explicitly **not** an employee-activity monitor.

2. Use cases by role

Employee — self-service (MVP)

Prompt	Tools	Returns
"What's on my plate this week?"	<code>list_my_assignments</code> , <code>get_work_log_summary</code>	Open tasks, due dates, what they closed
"Give me everything on the checkout task."	<code>get_task</code> , <code>list_task_notes</code>	Task detail + latest comments + linked milestone
"What am I blocked on?"	<code>search_tasks(status: blocked)</code>	Blockers with who/what they're waiting on

The trust win — the employee uses it on themselves, less status-reporting overhead, zero surveillance feel.

Team lead / manager (*MVP reads* → *Phase 2 writes*)

Prompt	Tools	Returns
"Summarize what my team shipped this week."	<code>list_team_assignments</code> , <code>get_work_log_summary</code> (prompt: <code>weekly_team_summary</code>)	Clean weekly rollup
"Who's overloaded right now?"	<code>get_team_capacity_summary</code>	Workload balance (aggregated)
"Run today's standup view."	prompt: <code>daily_standup_review</code>	Yesterday / today / blockers
"Prep me for my 1:1 with Kunal."	prompt: <code>one_on_one_prep</code> (scoped: report only)	Assignments, recent progress, blockers
Phase 2: "Note design is blocked on copy, move it to In Review."	<code>add_task_note</code> , <code>update_task_status</code>	Writes — confirmation required

PMO / project ops (*MVP*)

Prompt	Tools	Returns
"Which projects slipped this week?"	<code>list_project_milestones</code> , <code>search_tasks</code>	Milestones with no recent activity, flagged
"Health check on the Q3 launch."	<code>get_project</code> , prompt: <code>project_risk_review</code>	Status, stalled tasks, overloaded owners
"Which cross-project dependencies are stuck?"	<code>search_tasks(dependency, blocked)</code>	Blocked dependencies across projects

Org head / executive (*MVP — the command center*)

Prompt	Tools	Returns
"What's going on across every project?"	org-wide <code>get_project</code> + summaries	Portfolio status, one view

Prompt	Tools	Returns
"How much of Q3 is done, how much left?"	<code>get_project</code> , <code>list_project_milestones</code>	Completion % + remaining scope
"How many people are in today?"	<code>get_team_capacity_summary</code>	"92% present, 4 on leave" (aggregate)
"Top three risks across the company."	<code>project_risk_review</code> (portfolio)	Ranked risks with source projects

Org-wide scope is legitimate here — the org head's reach is the whole org. Inheritance gives them everything at the project/portfolio/capacity level.

HR / People (MVP — thresholded aggregates only)

Prompt	Tools	Returns
"How's team capacity trending this quarter?"	<code>get_team_capacity_summary</code> (cohort-thresholded)	Team-level signal, never raw activity

What KnowLens won't do: rank individuals by activity or output a per-person behavioral feed — the surveillance path, deliberately blocked for every role including leadership.

3. Top enterprise use cases (cross-role)

1. Conversational status & standup generation
2. Project risk / slippage detection
3. Workload & capacity balancing (aggregated)
4. Self-service "my week" briefings
5. Manager 1:1 prep packs
6. Portfolio rollups for PMO & leadership
7. (Phase 3) Scheduled briefings ("every Monday, org health summary")

4. Data scope & governance rules

Treatment	Applies to
Exposed (real-time)	Task status, assignments, project metadata, milestones
Cached	Org structure, project lists, slow-changing metadata

Treatment	Applies to
Summarized / aggregated	Capacity, engagement, productivity signals — computed deterministically, thresholded by cohort
Redacted	PII inside free-text notes/comments, before the model sees it
Hidden / never raw to LLM	Individual activity streams, person-level attendance, comp & sensitive PII, anything below the caller's permission scope

Three mechanics run under every query: (1) the agent only sees what the caller's role already grants; (2) people-signals return aggregated, never raw; (3) every call — who asked, which tool, what params, what returned — is logged.

5. Business model

Category: Internal AI enablement delivering a governed workforce- & project-intelligence layer — a blend of task intelligence + productivity intelligence. Not HR analytics, not monitoring.

Two intents

Intent	What it is	Model
Internal (build for Know AI)	Intelligence layer on your own KnowCore	Cost-center leverage; ROI = time + risk, not revenue
Product (sell to clients)	"Agent-ready workforce intelligence"	B2B; competes on governance depth, not breadth

ROI / value drivers: reporting-time collapse (standups, rollups, status lookups → seconds) · fewer status meetings · earlier risk detection → fewer slipped milestones · analytics without a BI project.

Monetization paths (if productized)

Model	Pros	Cons	Best when
Internal platform	Full control; data never leaves; simplest governance	Pure cost center	Know AI's own use
Self-hosted product	Sells to regulated / data-sovereignty buyers	High support burden	Healthcare / legal / gov / EU
Enterprise SaaS	Recurring revenue, fast iteration	You hold sensitive data → heavy compliance; keenest buyers can't use SaaS	SMB / non-regulated
Open-core	Community adoption + monetize enterprise governance (SSO, audit, RBAC depth)	License discipline, fork risk	Both adoption + revenue

The tension: the data-sovereignty angle that makes self-hosted/open-core attractive is the same reason pure SaaS is a hard sell for workforce data. If productizing, open-core fits how this category gets bought. Apache-2.0 components keep that door open.

6. The defining line

Org-wide scope: yes. Raw per-person activity as a monitoring feed: no.

"Project X is 60% done, team at 110% capacity, 4 on leave" is operations intelligence — everything leadership actually needs to know the org is on track. A continuous individual-level activity stream piped to an LLM is a different thing: legally exposed, trust-destroying, the weakest-adopted category in the market, and it doesn't even serve the org head's goal. Everything aggregates up; people-data resolves down only as far as the legitimate management chain already allows.

7. Phased value

Phase	Identity	Use cases unlocked
MVP	Read-only analytics assistant	All reads above — self-service, team rollups, PMO/exec views, HR aggregates
Phase 2	Workflow assistant	Scoped, confirmed writes — notes, status transitions, follow-up tasks
Phase 3	Project-intelligence copilot	Pre-computed risk intelligence + proactive scheduled briefings (needs an orchestration layer — MCP is pull-only)

Build first: the read-only MVP. Lowest risk, fastest ROI, and it forces the governance spine (auth, audit, row-level inheritance) that everything else depends on.