

Your close rate isn't a pitch problem.

It's a sequencing problem. The transcript shows you exactly where.

A field guide for reviewing a week of sales calls in one session and isolating the specific phrases, timing habits, and conversation patterns that are costing you closes.

Batch size	Audit time	Weekly cadence
5-20 calls	One session	6-10 minutes

Figures are based on the system as described. Actual time will vary by call volume and transcript quality.

Most sellers diagnose the wrong variable.

When close rate drops, the first instinct is to adjust price. The second is to rework the pitch. Both feel productive. Neither addresses the real pattern.

In most cases, the issue is not what you said. It is when you said it. Price comes up before the prospect has named a problem worth solving. Silence gets filled with a concession. The pitch starts before the diagnosis is finished. These are timing failures, not content failures — and they are invisible in memory because memory compresses the sequence into a feeling: *that call went well* or *they weren't ready to buy*.

Transcripts do not compress. They preserve the minute-by-minute record: the exact line where the prospect went quiet, the exact word that introduced price, the pause that got filled with a discount instead of a question. That is where the close rate leak is.

The difficulty has always been volume. Replaying calls one by one, at full length, takes most of a working day. Sellers who do it at all do it once, draw one conclusion, and move on. The pattern — the thing that is happening across every low-scoring call, not just the worst one — never surfaces.

This system uses Claude to compress that process. You load a batch of transcripts, run a single coaching prompt, and receive a per-call diagnostic with exact phrase flags and a directional close-rate score for each call. The value is not in any one call's output. It is in what the batch reveals when every call is scored against the same framework at the same time.

The system has four steps. Run them in order the first time. After that, Step 4 maintains itself.

The pattern this system consistently surfaces: Sellers who improve close rate are not better at closing. They are better at sequencing — specifically, they move the diagnosis earlier and the price conversation later. The pitch is rarely the problem. The order is.

Step 01 • Load the Batch

Start with transcripts, not with memory.

Sales coaching built from memory almost always reflects the calls you remember most, not the calls that are most representative. The call where the prospect pushed hard on price stays vivid. The call where you pitched too early and the prospect went politely quiet is already fading. Memory selects for emotional intensity, not diagnostic accuracy.

Transcripts are the corrective. They record what was said, by whom, and when — with enough fidelity that a coaching model can identify behavioral patterns across a batch of calls.

The practical setup: Most major call recording platforms (Zoom, Gong, Chorus, Fathom, Fireflies) export transcripts as .txt or .vtt files. In Zoom: *Recordings* → *Download* → *Transcript*. Export the last 10–20 calls and drop them into a single Claude conversation. Do not filter before you load. Include the calls you felt good about and the ones you would rather not revisit. The pattern lives in the contrast.

On batch size: The full system works best with 15–20 calls — enough volume to distinguish recurring habits from one-off mistakes. If you are running this for the first time, 5 calls is a workable starting point; you will see a first signal, though not a durable pattern. Build toward 15–20 over successive weeks.

Prompt to run:

"I'm dropping [X] sales call transcripts from the past 30 days. Read them in order. Do not analyze yet. Confirm you have all [X] and flag any call under 4 minutes — those are unlikely to be full sales conversations."

Simulated output:

```
✓ 18 transcripts loaded, timestamps preserved
✓ Speaker labels detected: Rep / Prospect
✓ 2 calls flagged under 4 min (calls 7, 15)
```

```
— batch summary —
Usable calls      18
Avg duration      24 min
Shortest (flagged) 3 min
Longest           51 min
Speaker ratio     Rep 61% · Prospect 39%

Ready for coaching prompt
```

Note: This output is illustrative. Your platform, call length, and speaker labels will differ.

What this step produces

- A clean batch of timestamped transcripts ready for pattern analysis
- Early flags on calls too short to be diagnostically useful
- A baseline speaker ratio — useful context before coaching begins

Where it breaks

- Audio-only recordings require transcription before loading (Whisper, Otter, or your platform's built-in transcription)
 - Transcripts with only one speaker captured produce weaker coaching output; Claude cannot analyze what the prospect said if it was not recorded
 - Batches under 8 usable calls are thin for pattern detection — flag and note it rather than drawing firm conclusions
-

Step 02 • Run the Diagnostic

Ask for a coach, not a summary.

Most people prompt Claude to "summarize the call." A summary tells you what happened. It does not tell you what to change. The framing that produces actionable coaching is different: you are not asking for a recap. You are asking for a behavioral audit — where did the conversation lose momentum, what caused it, and what should have happened instead.

Claude scores each call on a 0–100 directional scale. The score reflects structural and linguistic patterns in the call — how early the pitch arrived, whether diagnosis questions preceded the price conversation, how silence was handled. It is not a verified performance

metric and should not be read as one. Its value is in the spread: which calls clustered around 30, which clustered around 70, and what distinguishes them.

Run this prompt across the full usable batch in a single pass. Do not ask for one call at a time. The batch framing is what allows Claude to identify patterns across calls rather than analyzing each one in isolation.

Prompt to run:

"Act as my sales coach. Review each call individually. For each one, identify: 1. Where momentum dropped — give me the minute and the exact line. 2. The specific phrase I used that caused it. 3. What I should have said instead. Then give each call a close-rate score from 0-100. Note that this score is directional — I want to understand the spread, not benchmark myself against an external standard."

Simulated output:

```
✓ 18 calls reviewed
✓ Momentum drops identified per call
✓ Directional close-rate scores assigned

- call 01 · score 72 -----
Momentum drop  Min 9 - "let me know what budget works"
Phrase          "we can be flexible on pricing"
Fix            Hold silence for 8 seconds. Ask:
                "What makes this the right time for you?"

- call 04 · score 31 -----
Momentum drop  Min 6 - pitch started before pain named
Phrase          "so here's what we do..."
Fix            Ask 3 diagnosis questions first.
                Pitch only after the problem is stated.

- call 11 · score 58 -----
Momentum drop  Min 19 - "I can do 15% off if you sign this week"
Phrase          Discount filled the silence
Fix            Silence is not a signal to concede.
                Wait 10 seconds. Ask: "What would make this feel right?"
```

Scores and outputs are illustrative. Your actual results will depend on call structure and transcript quality.

What this step produces

- A per-call momentum map with timestamps and exact phrases
- A directional score spread that shows which calls were structurally stronger
- Replacement language grounded in your actual calls, not generic scripts

Where it breaks

- The score reflects call structure, not verified outcome. A high score on a lost deal means the call was solid; something outside the conversation closed it. Do not conflate score with result.
 - One-sided transcripts (prospect audio missing) produce coaching that is weaker on diagnosis questions and prospect engagement
 - The scores are a diagnostic tool, not a performance review. Attaching them to self-assessment rather than behavior change is a misuse
-

Step 03 • Extract the Pattern

One bad call is a bad day. The same thing across nine calls is a habit.

After the per-call diagnostic is complete, run one more prompt. This one does not look at individual calls. It looks at what every low-scoring call has in common.

This is where the real coaching lives — not in the specific call that went wrong, but in the behavior that recurs across every call that went wrong. The distinction matters because habits require rules, not reminders. Deciding to "listen better" or "be more patient on price" does not hold under the pressure of a live call. A specific, testable rule does: *no price before minute 15, three diagnosis questions before the pitch, ten seconds of silence before any concession*. Rules have a pass/fail condition you can check after every call. Resolutions do not.

Cap the output at three patterns. More than three cannot be held in working memory during an active call.

Prompt to run:

"Look at the calls scored under 50. What do they have in common — the same timing pattern, the same phrase type, the same speaker ratio? Give me the 3 recurring behaviors that appear across the most calls. For each one, write a specific, testable rule I can apply starting next Monday."

Simulated output:

- ✓ 9 calls below 50 analyzed
- ✓ 3 recurring behaviors isolated

— pattern 01 · timing —————

Observation Pitch started avg at min 6.2
 Prospect silent by min 9 in 7 of 9 calls
Rule No pitch before 3 diagnosis questions answered

— pattern 02 · silence —————

Observation Avg pause before rep speaks: 4 seconds
 Discount or concession followed in 8 of 9 calls
Rule Wait 10 full seconds before speaking after prospect pause
 Ask: "What would make this feel right?"

— pattern 03 · ratio —————

Observation Rep: 61% of call · Prospect: 39%
Target Rep: 40% · Prospect: 60%
Rule After each prospect answer, ask one follow-up
 before continuing

What this step produces

- Three testable rules derived from your actual calls, not generic advice
- A behavioral root cause, not a surface-level symptom
- A specific set of changes to take into your next call

Where it breaks

- Patterns from one 15–20 call batch should be treated as hypotheses, not conclusions. Validate against the next 10 calls before treating them as durable
- If the low-scoring half contains fewer than 6 calls, the pattern is too thin to rely on. Run the step again after a fuller batch
- Three rules is the cap. Adding a fourth dilutes your attention on all of them

Step 04 · Run the Weekly Audit

One batch is a diagnosis. A weekly rhythm is how the diagnosis becomes improvement.

Running this system once tells you what is currently broken. Running it every week tells you whether the fix is holding — and where the next leak has appeared.

The weekly audit is the same system with a shorter batch (the week's calls, typically 3–5) and one additional prompt: compare this week's scores to last week's. Did pitch timing

improve? Is the silence pattern holding? Did a new habit appear? The output is not a full re-analysis. It is a delta: one thing that improved, one thing that did not, and one specific focus for the following week.

The Friday cadence works because it closes the feedback loop before the week's pattern resets. The goal is not a perfect week. It is a measurable week.

Prompt to run:

"Here are this week's transcripts. Score each call using the same framework. Then compare: did pitch timing improve from last week's audit? Did silence tolerance improve? Show me the one metric that improved and the one that did not. Give me a single focus rule for next week."

Simulated output:

```
✓ 4 calls loaded, scored, compared to prior week

— this week · avg score 58 —————
Last week avg   44 → This week avg 58

— improved —————
Pitch timing     Avg min 10.4 (from 6.2)
Diagnosis Qs     Avg 2.8 per call (from 1.2)

— still breaking —————
Silence gap      Avg 5 sec before rep speaks (was 4 sec)
Discount habit   Appeared in 2 of 4 calls

— next week rule —————
Focus            Silence tolerance
Target           10 full seconds before speaking
Track            Note the exact minute you broke the rule
```

What this step produces

- A week-on-week score comparison that makes improvement visible rather than merely felt
- A single focus rule each Friday — specific enough to test on Monday's first call
- A compounding data set that, over four weeks, shows a genuine trend rather than a single data point

- This step requires consistency. Skipping a Friday breaks the delta comparison; the value of the system is in continuity, not any single run
 - Weeks with fewer than 3 calls produce thin data. Note it; do not draw firm conclusions from a single call
 - The audit finds patterns. It does not change them. The rule you write on Friday must carry into the calls the following week — the system creates clarity, not automatic improvement
-

What sellers who use this system report

The following observations are drawn from sellers who have run this audit over four or more consecutive weeks. They are representative patterns, not controlled outcomes.

On timing: The most common discovery is that price comes up earlier than the seller believed — often before the prospect has described the problem in their own words. Adjusting pitch timing is typically the first change that produces a visible score improvement.

On silence: Most sellers discover they are filling silence within 4–6 seconds, usually with a concession or a price qualifier. Extending that gap to 10 seconds is uncomfortable in practice. Sellers who hold it report that prospects fill the silence with more useful information than any follow-up question would have generated.

On speaker ratio: A 60/40 split (prospect to rep) appears consistently in higher-scoring calls across multiple users of this system. Reaching that ratio does not require talking less — it requires asking questions that generate longer answers.

On the pattern that most surprised sellers: The issue is almost never the pitch quality. It is the sequence. When diagnosis questions move earlier and price moves later, the pitch lands in a different context — one where the prospect has already named the problem — and the conversation dynamic shifts without the pitch itself changing.

The system at a glance

Step	What it analyzes	Input	Output
01 · Load	Raw transcript data	5–20 .txt or .vtt files	Clean batch, flagged and ready
02 · Diagnose	Per-call behavior	One prompt	Score + exact phrase flag per call
03 · Pattern	Cross-call habits	Low-scoring half	3 root patterns + 3 testable rules
04 · Automate	Week-on-week delta	Weekly transcripts	One-page trend report + next rule

Run order matters. The pattern report requires per-call scores. The weekly automation requires an established baseline and working rules. Each step gates the next. Do not skip to Step 03 before every call has a score.

How to start

The full system uses 15–20 calls to surface durable patterns. You do not need 20 calls to begin.

Take the last 3–5 calls you ran. Export the transcripts from your recording platform. Drop them into Claude as separate files or pasted text. Run the Step 02 coaching prompt.

When the output comes back, read it for one thing only: the phrase that appears in the lowest-scoring call that also shows up in at least one other. That is your first pattern. That is what you work on this week.

After four consecutive Friday audits, you will have enough data to distinguish durable habits from weekly noise. That is when the system starts to compound.

Run this prompt now

Copy this into Claude with your transcripts loaded:

"Act as my sales coach. Review each call individually. For each one, tell me: where momentum dropped (give me the minute and the exact line), the specific phrase I used that caused it, and what I should have said instead. Score each call 0–100 on a directional close-rate scale. Flag any phrase that appears in more than one call — those are habits, not one-off mistakes."

One session. Your calls. Your exact words.

That is what changes the pattern.

Close-Rate Audit · A Field Guide Transcripts · Diagnosis · Pattern · Automation