

JYOTHIR RAGHAVALU BHOGI

[Jyothir-dev | LinkedIn](#)
[Jyothir-dev \(Jyothir\) \(github.com\)](#)

Email: jyothirraghavalu369@gmail.com
Mobile: +91 6281073800

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL
- **AI/ML:** Machine Learning, Deep Learning, PyTorch, TensorFlow (Basic), Scikit-learn, Model Evaluation
- **Generative AI:** LLMs, RAG, Hugging Face, Embeddings, Vector Search
- **Data:** NumPy, Pandas, Data Preprocessing, Data Validation, ETL
- **Backend:** FastAPI, REST APIs
- **Deployment:** Docker, AWS, ML Deployment
- **Engineering:** Git, GitHub, Testing, Debugging

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)

June 2026 - Present

 - Built Generative AI workflows using LLM inference, prompt engineering, context engineering, and evaluation.
 - Developed Python backend services, REST APIs, and AI pipelines for LLM, search, and automation workflows.
 - Designed RAG pipelines using embeddings and vector search to improve contextual retrieval and AI reliability.
 - Developed AI evaluation pipelines with regression testing, error analysis, validation, and performance monitoring.
 - Applied LangGraph agentic workflows for task orchestration and multi-step LLM applications.
- Data Engineer Intern | flickDONE | [LINK](#)

June 2025 - August 2025

 - Built Python ML pipelines using NumPy and Pandas, processing 50,000+ records for training and validation.
 - Engineered automated preprocessing and data validation workflows, reducing manual verification by 70%.
 - Optimized AWS ML pipelines, improving execution speed by 25% through bottleneck analysis.
 - Developed reusable AI workflows for conversational AI, document processing, and business automation.
 - Applied Docker and CI/CD for containerized ML deployment, improving reliability and scalability.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)

April 2026 - June 2026

 - Built a RAG application using Python, PyTorch, Hugging Face, embeddings, and vector search.
 - Designed asynchronous RAG retrieval architecture, improving semantic search efficiency by 50%.
 - Developed modular retrieval pipelines with context management, reducing end-to-end inference latency by 50%.
 - Improved AI output reliability through model evaluation, grounding validation, testing, debugging, and response quality checks.
- AI Search & Recommendation Platform | [LINK](#)

June 2025 - August 2025

 - Designed an AI-powered semantic search platform for retrieval and recommendation workflows.
 - Built FastAPI REST APIs for ML inference, model serving, semantic search, and AI workflows.
 - Integrated Databricks infrastructure, reducing AI routing latency by **35%**.
 - Containerized services with Docker, improving deployment reliability and scalability by **25%**.
- Morning Market Brief | [LINK](#)

June 2024 - August 2024

 - Built Generative AI workflows using OpenAI APIs, LLM inference, and prompt engineering for financial summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by 45%.
 - Implemented prompt validation, LLM evaluation, and quality checks, improving model reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)

August 2024

 - Built ML models using Python, NumPy, Pandas, and PyTorch for development and evaluation.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with Large Language Models (Coursera) | [CERTIFICATE](#)

March 2024

 - Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows, AI output evaluation, and RAG pipeline performance validation.