

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL
- **Generative AI:** LLMs, RAG, LangChain, LangGraph, Prompt Engineering, Embeddings, Agentic AI, Prompt Design, LLM Workflows, Structured Output
- **RAG & Retrieval:** Vector Databases, FAISS, Semantic Search, Chunking, Indexing, Retrieval Evaluation
- **AI Engineering:** FastAPI, REST APIs, Tool Calling, API Development, Microservices, Performance Optimization
- **Production:** AWS, Docker, Git, CI/CD, Automated Testing, Scalable AI Systems

WORK EXPERIENCE

AI Intern | Know AI | [LINK](#)
June 2026 - August 2026

- Built AI and LLM applications using Python, NLP, model inference, experimentation, and model evaluation.
- Developed Python AI services and REST APIs integrating LLMs, prompt workflows, tool calling, search, and automation.
- Designed **RAG pipelines** using embeddings, vector databases, and semantic search for knowledge-grounded retrieval.
- Developed AI/ML evaluation pipelines for model testing, error analysis, output validation, and reliability analysis.
- Improved AI workflow reliability through evaluation, monitoring, debugging, performance tuning, and latency optimization.

Data Engineer Intern | flickDONE | [LINK](#)
June 2025 - August 2025

- Built Python-based ML data pipelines processing 50,000+ records for preprocessing, feature engineering, and validation.
- Engineered automated **data preprocessing and validation workflows**, reducing manual verification by 70%.
- Optimized AWS ML pipelines, improving execution speed by 25% through bottleneck analysis and optimization.
- Developed reusable AI workflows for **document processing, LLM applications, and enterprise automation**.
- Collaborated with engineering teams to improve **scalability, reliability, maintainability, and deployment** of ML platforms.

PROJECTS

Ledgerly — Multi-Tenant SaaS Billing Platform | [LIVE](#)
April 2026 - September 2026

- Built a **multi-tenant SaaS billing platform** using Next.js, NestJS, PostgreSQL, Prisma, and JWT authentication.
- Implemented **tenant-isolated customer and invoice workflows**, dashboard balances, and OWNER/STAFF RBAC.
- Built **Stripe payment workflows and webhook processing** with idempotency, transactional updates, and notifications.
- Added **rate limiting, health checks, secure error handling, and responsive UI** for production-oriented reliability.
- Developed 174 automated tests across 21 suites, covering authentication, billing, payments, notifications, and tenant isolation.

LangGraph Workflow Orchestrator | [LIVE](#)
August 2025

- Built and deployed a stateful **Agentic AI platform** using Python, LangGraph, LangChain, FastAPI, and Gemini.
- Designed multi-step LLM conversation workflows with state management, conditional routing, and error recovery.
- Developed modular components for **task decomposition, tool calling, function calling, and agent coordination**.
- Exposed AI workflows through **FastAPI REST APIs**, enabling scalable AI task automation.
- Improved workflow reliability through **checkpointing, testing, failure handling, and monitoring**.

Production-Grade RAG Engine | [LIVE](#)
June 2024 - August 2024

- Built an end-to-end **RAG application** using Python, FastAPI, LangChain, Qdrant, and AWS.
- Developed document preprocessing, chunking, embeddings, indexing, retrieval, and response generation pipelines.
- Implemented **semantic search** with Qdrant for efficient context retrieval.
- Added reranking and retrieval evaluation to improve retrieval quality, context relevance, and AI response performance.
- Improved AI service **reliability and maintainability** through testing, monitoring, fallback execution, and error handling.

CERTIFICATES

Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)
August 2024

- Built machine learning models using PyTorch and applied ML techniques for model development and evaluation.
- Developed **Dockerized environments** supporting reproducible experimentation and deployment

Generative AI with Large Language Models (Coursera) | [CERTIFICATE](#)
March 2024

- Applied prompt engineering techniques for production LLM applications.
- Built **retrieval workflows, AI output evaluation, and RAG pipeline** performance validation.