

1 ☐ JYOTHIR RAGHAVALU BHOGI
2 ☐
3 ☐ Generative AI Engineer | LLMs, RAG & AI Agents | Python, LangChain,
4 ☐ LangGraph, FastAPI, AWS
5 ☐
6 ☐ jyothirraghavalu369@gmail.com | +91 6281073800 | LinkedIn | GitHub
7 ☐
8 ☐ Generative AI Engineer building production-oriented LLM, RAG, and
9 ☐ agentic AI systems – RAG pipelines with embeddings and vector search,
10 ☐ multi-step agent orchestration, LLM evaluation, and the backend APIs
11 that expose all of it. B.Tech CSE candidate (2026) with two AI/ML
12 internships and three shipped GenAI projects behind REST APIs.

14 EDUCATION

16 Lovely Professional University Jalandhar, India
18 Bachelor of Technology, Computer Science & Engineering July 2022 - May
19 2026

21 SKILLS SUMMARY

23 Generative AI: LLMs, RAG, Prompt Engineering, Embeddings, Vector Search,
24 Semantic Search
26 AI Agents / Orchestration: LangChain, LangGraph, Agentic Workflows, Task
27 Decomposition, Checkpointing
29 AI/ML: Machine Learning, Deep Learning, ML Pipelines, Model Inference,
30 Model Evaluation, PyTorch, Hugging Face
32 APIs & Deployment: FastAPI, REST APIs, Model Serving, AWS, Docker,
33 CI/CD, Model Deployment
35 Data: Python, SQL, PySpark, Apache Spark, Data Preprocessing, Feature
36 Engineering, ETL, Batch Processing
38 Engineering: Performance Optimization, Debugging, Workflow Automation,
39 Git, GitHub

41 WORK EXPERIENCE

43 Generative AI Engineer Intern | Know AI June 2026 - August 2026
45 - Built Generative AI workflows using LLM inference, NLP, prompt
46 engineering, and structured model evaluation.
48 - Designed RAG pipelines using embeddings and vector search to improve

- 49 contextual retrieval and AI reliability.
- 50
- 51 - Built backend services and REST APIs (FastAPI) integrating LLMs, AI
- 52 inference, search, and automation workflows.
- 53
- 54 - Developed AI/ML evaluation pipelines for model testing, error
- 55 analysis, output validation, and performance analysis.
- 56
- 57 - Conducted model evaluation, error analysis, and output validation
- 58 that directly improved ML reliability.
- 59

60 Data Engineer Intern | flickDONE June 2025 - August 2025

61

- 62 - Built Python-based ML data pipelines processing 50,000+ records for
- 63 preprocessing, training, and validation.
- 64
- 65 - Engineered automated preprocessing and data validation workflows,
- 66 reducing manual verification by 70%.
- 67
- 68 - Optimized AWS ML pipelines, improving execution speed by 25% through
- 69 bottleneck analysis.
- 70
- 71 - Developed reusable AI workflows for conversational AI, document
- 72 processing, and business automation.
- 73
- 74 - Collaborated with engineering teams to improve scalability,
- 75 reliability, and maintainability of ML platforms.
- 76

77 PROJECTS

78

79 Production-Grade RAG Engine June 2024 - August 2024

80

81 Python, FastAPI, LangChain, Qdrant, AWS

82

- 83 - Built an end-to-end production RAG pipeline: document ingestion,
- 84 chunking, embedding, vector retrieval, and context generation.
- 85
- 86 - Implemented semantic search with a Qdrant vector database for
- 87 efficient document context retrieval.
- 88
- 89 - Added reranking and retrieval evaluation to improve context
- 90 relevance and response quality.
- 91
- 92 - Improved system reliability through fallback execution, testing,
- 93 monitoring, and error handling.
- 94

95 LangGraph Workflow Orchestrator August 2025

96

97 Python, LangGraph, LangChain, FastAPI, Gemini
98
99 - Built a stateful AI agent orchestration platform with multi-step
100 workflows, state management, and conditional routing.
101
102 - Developed modular components for task decomposition, tool execution,
103 and automated agent coordination.
104
105 - Exposed workflow execution through FastAPI REST APIs, enabling
106 scalable, reusable AI task automation.
107
108 - Improved workflow reliability through checkpointing, failure
109 handling, testing, and execution monitoring.

110
111 Regulatory Compliance RAG Assistant April 2026 - June 2026

112
113 Python, Streamlit, LangChain, Embeddings, Vector Search

114
115 - Built a Generative AI compliance assistant with document ingestion
116 and semantic retrieval pipelines.
117
118 - Implemented context-grounded generation to deliver relevant and
119 reliable compliance responses.
120
121 - Developed interactive document upload and querying workflows using
122 Streamlit.
123
124 - Improved response reliability through retrieval testing, evaluation,
125 debugging, and quality checks.

126
127 CERTIFICATES

128
129 Generative AI with Large Language Models – Coursera March 2024

130
131 - Applied prompt engineering techniques for production LLM
132 applications.
133
134 - Built retrieval workflows, AI output evaluation, and RAG pipeline
135 performance validation.

136
137 Python, Data Science & Machine Learning – CipherSchools August 2024

138
139 - Built machine learning models using PyTorch and applied ML
140 techniques for model development and evaluation.
141
142 - Developed Dockerized environments supporting reproducible
143 experimentation and deployment.

144