

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL, Object-Oriented Programming (OOP)
- **AI/ML:** LLM Applications, AI Agents, Agentic Workflows, Prompt Engineering, Model Evaluation, Performance Optimization
- **RAG:** Embeddings, Vector Databases, Retrieval-Augmented Generation
- **Frameworks:** PyTorch, Hugging Face, Transformers
- **Deployment:** REST APIs, ML Inference, Model Serving, Docker, AWS, API Integration
- **Engineering:** Backend Services, Workflow Automation, System Debugging, Git, GitHub

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)** June 2026- Present
- Developed production AI applications using Python, PyTorch, Hugging Face, Transformers, and LLM-based workflows.
 - Built backend services and REST APIs integrating AI models, ML inference, and automation pipelines.
 - Designed RAG pipelines integrating embeddings, vector databases, and retrieval systems to improve AI reliability.
 - Built scalable AI workflows using LLM inference, prompt engineering, evaluation pipelines, and automation techniques.
 - Conducted model evaluation, error analysis, AI output validation, and performance optimization to improve ML reliability.
- Data Engineer Intern | flickDONE | [LINK](#)** June 2025- August 2025
- Built scalable ML data pipelines processing 50,000+ samples for training, validation, and text classification workflows.
 - Engineered automated preprocessing and tokenization workflows, reducing manual verification by **70%**.
 - Optimized AWS-based ML pipelines and data workflows, improving execution speed by 25% through bottleneck analysis.
 - Developed reusable AI workflows for conversational AI, document processing, and business automation.
 - Collaborated with engineering teams to improve scalability, reliability, and maintainability of ML platforms.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)** April 26- June 2026
- Built an end-to-end AI pipeline using Python, PyTorch, Hugging Face, embeddings, RAG, and LLM workflows.
 - Designed asynchronous retrieval architecture, improving semantic search efficiency by 50%.
 - Developed modular retrieval and inference pipelines, reducing end-to-end latency by 50%.
 - Improved zero-shot fraud detection accuracy by 30% through model evaluation and AI output validation.
- AI Search & Recommendation Platform | [LINK](#)** June 25- August 2025
- Designed an AI-powered semantic search platform for retrieval and recommendation workflows.
 - Built scalable FastAPI REST APIs for ML inference and backend orchestration workflows.
 - Integrated Databricks infrastructure, reducing AI routing latency by **35%**.
 - Containerized services with Docker, improving deployment reliability and scalability by **25%**.
- Morning Market Brief | [LINK](#)** June 24- August 2024
- Built AI workflows using OpenAI APIs, LLM pipelines, and prompt engineering for financial summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by 45%.
 - Implemented prompt validation, LLM evaluation, and quality checks, improving model reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)** August 2024
- Built ML models using PyTorch.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with LLMs (Coursera) | [CERTIFICATE](#)** March 2024
- Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows, AI output evaluation, and RAG pipeline performance validation.