

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL, PySpark
- **AI/ML:** Machine Learning, Deep Learning, ML Pipelines, Model Inference, Model Evaluation, PyTorch, Hugging Face
- **Generative AI:** LLMs, Generative AI, RAG, Prompt Engineering, Embeddings, Vector Search, Semantic Search
- **Data:** Apache Spark, PySpark, Data Preprocessing, Feature Engineering, ETL, Batch Processing
- **APIs:** FastAPI, REST APIs, HTTP APIs, Model Serving
- **Deployment:** AWS, Docker, CI/CD
- **Engineering:** Performance Optimization, Scalability, Workflow Automation, Git, GitHub

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)June 2026 - Present
 - Built Generative AI workflows using LLM inference, NLP, prompt engineering, and model evaluation.
 - Built backend services and REST APIs integrating LLMs, AI inference, search, and automation workflows.
 - Designed RAG pipelines using embeddings and vector search to improve contextual retrieval and AI reliability.
 - Developed AI/ML evaluation pipelines for testing, error analysis, AI output validation, and performance analysis.
 - Conducted model evaluation, error analysis, AI output validation, and model optimization to improve ML reliability.
- Data Engineer Intern | flickDONE | [LINK](#)June 2025 - August 2025
 - Built Python-based ML data pipelines processing 50,000+ records for preprocessing, training, and validation.
 - Engineered automated preprocessing and tokenization workflows, reducing manual verification by **70%**.
 - Optimized AWS ML pipelines, improving execution speed by 25% through bottleneck analysis.
 - Developed reusable AI workflows for conversational AI, document processing, and business automation.
 - Collaborated with engineering teams to improve scalability, reliability, and maintainability of ML platforms.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)April 2026 - June 2026
 - Built a Generative AI and NLP RAG pipeline using Python, PyTorch, Hugging Face, embeddings, vector search, and LLMs.
 - Designed asynchronous RAG retrieval architecture, improving semantic search efficiency by 50%.
 - Developed modular retrieval and inference pipelines, reducing end-to-end inference latency by 50%.
 - Improved AI output reliability through model evaluation, prompt validation, and response quality checks.
- AI Search & Recommendation Platform | [LINK](#)June 2025 - August 2025
 - Designed an AI-powered semantic search platform for retrieval and recommendation workflows.
 - Built scalable FastAPI REST APIs for ML inference, HTTP-based model serving, semantic search, and backend orchestration.
 - Integrated Databricks infrastructure, reducing AI routing latency by **35%**.
 - Containerized services with Docker, improving deployment reliability and scalability by **25%**.
- Morning Market Brief | [LINK](#)June 2024 - August 2024
 - Built Generative AI workflows using OpenAI APIs, LLM inference, and prompt engineering for financial summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by 45%.
 - Implemented prompt validation, LLM evaluation, and quality checks, improving model reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)August 2024
 - Built machine learning models using PyTorch and applied ML techniques for model development and evaluation.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with Large Language Models (Coursera) | [CERTIFICATE](#)March 2024
 - Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows, AI output evaluation, and RAG pipeline performance validation.