

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL
- **Data:** Pandas, Vector Databases (ChromaDB)
- **AI:** Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Prompt Engineering, Agentic AI, Tool Calling
- **Frameworks:** PyTorch, Hugging Face, Transformers, LangChain, LlamaIndex
- **Backend:** FastAPI, REST APIs, Model Inference, Model Deployment, Docker
- **Evaluation:** Model Evaluation, AI Output Validation, Performance Optimization
- **Engineering:** Debugging, Git, GitHub

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)

June 2026- Present

 - Developed production AI applications using Python, PyTorch, Hugging Face, Transformers, and LLMs.
 - Built backend services and REST APIs supporting LLM inference, scalable AI workflows, and deployment.
 - Designed modular RAG pipelines integrating vector databases to improve retrieval quality and response consistency.
 - Built agentic AI workflows using prompt engineering, tool calling, and AI orchestration for multi-step LLM tasks.
 - Conducted model evaluation, AI output validation, debugging, and performance optimization to improve AI reliability.
- Data Engineer Intern | flickDONE | [LINK](#)

June 2025- August 2025

 - Built scalable NLP pipelines processing 50,000+ datasets for parallel text classification.
 - Engineered automated preprocessing and tokenization workflows, reducing manual verification by **70%**.
 - Optimized AWS-based ML pipelines, improving execution speed by **25%** through bottleneck analysis.
 - Developed reusable AI workflows supporting conversational AI, intelligent document processing, and workflow automation.
 - Collaborated with product and engineering teams to improve scalability, reliability, and maintainability of production AI systems.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)

April 26- June 2026

 - Built a RAG system using Python, Hugging Face, ChromaDB, and vector retrieval pipelines.
 - Designed asynchronous retrieval architecture, improving semantic search efficiency by 50%.
 - Developed modular retrieval and inference pipelines, reducing end-to-end latency by 50%.
 - Improved zero-shot fraud detection accuracy by 30% through model evaluation and AI output validation.
- AI Search & Recommendation Platform | [LINK](#)

June 25- August 2025

 - Designed an LLM-powered semantic search platform supporting intelligent retrieval and recommendations.
 - Built scalable FastAPI REST APIs for AI inference and orchestration workflows.
 - Integrated Databricks infrastructure, reducing AI routing latency by **35%**.
 - Containerized services with Docker, improving deployment reliability and scalability by **25%**.
- Morning Market Brief | [LINK](#)

June 24- August 2024

 - Built automated AI workflows integrating OpenAI APIs for financial news summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by 45%.
 - Implemented prompt validation and quality checks, improving LLM reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)

August 2024

 - Built ML models using PyTorch.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with LLMs (Coursera) | [CERTIFICATE](#)

March 2024

 - Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows, AI output evaluation, and RAG pipeline performance validation.