

Jyothir Raghavalu Bhogi

+91 6281073800 | jyothirraghavalu369@gmail.com | linkedin.com/in/jyothir-dev | github.com/jyothir-369 | leetcode.com/u/jyothir-369/ | Immediate Joiner | Open to Relocation | Remote

SUMMARY

AI/ML Engineer building **production-oriented LLM, RAG, and GenAI applications** using Python, LangGraph, LangChain, FastAPI, PyTorch, scikit-learn, Gemini, and Hugging Face. Experienced in developing **end-to-end AI pipelines, embeddings, vector retrieval, evaluation, APIs, and deployment**, with hands-on work across data ingestion, model inference, latency optimization, reliability, monitoring, Docker, cloud, and CI/CD.

EDUCATION

Lovely Professional University

B.Tech in Computer Science and Engineering

Jalandhar, Punjab

Aug 2022 – May 2026

- **Relevant Coursework:** Artificial Intelligence, Machine Learning, NLP, Data Structures & Algorithms, Distributed Systems, DBMS, Statistics

EXPERIENCE

AI Intern

Know AI

June 2026 – Present

Ahmedabad, Gujarat

- Build **LLM, RAG, and GenAI applications** using Python, LangGraph, LangChain, FastAPI, Gemini, and Hugging Face.
- Develop **end-to-end AI workflows** covering data ingestion, retrieval, tool calling, structured outputs, evaluation, and backend integration.
- Implement **prompt engineering, LLM evaluation, regression testing, grounding checks, logging, and monitoring** to improve accuracy and reliability.
- Integrate AI systems with **REST APIs, vector databases, structured data, cloud services, and production-oriented backend pipelines**.

Data Engineer Intern

flickDone

June 2025 – Aug 2025

Remote

- Built **Python and SQL data pipelines** processing **50,000+ records** for analytics and machine learning workflows.
- Developed data **ingestion, cleaning, preprocessing, validation, and transformation** workflows using Pandas, NumPy, Python, and SQL.
- Improved **system reliability by 30%** and reduced **query latency by 25%** through data and query optimization.
- Implemented **Docker, GitHub Actions, and CI/CD** for reproducible testing, deployment, and production workflows.

PROJECTS

LangGraph Workflow Orchestrator | *Live Demo* | *Python, LangGraph, LangChain, Gemini, Chainlit* Aug 2026

- Built a **production-oriented Agentic AI workflow system** using LangGraph for stateful agent execution, task routing, and workflow orchestration.
- Developed reusable **agent, node, and workflow abstractions** with structured state management and tool-based execution.
- Implemented **testing, error handling, monitoring, and workflow validation** across approval, conditional, iterative, and parallel patterns.

Production-Grade RAG Engine | *Live Demo* | *Python, FastAPI, Gemini, Qdrant, Inngest* Jun 2026 – Aug 2026

- Built a **production-oriented RAG pipeline** for document ingestion, chunking, Gemini embeddings, vector retrieval, ranking, and grounded generation.

- Developed **FastAPI REST APIs and Streamlit UI** with source-aware retrieval, semantic search, and multi-provider configuration.
- Implemented **grounding controls, retrieval evaluation, fallback execution, and testing** to improve RAG accuracy and response reliability.

Regulatory Compliance RAG Assistant | *Live Demo* | *Python, LangChain, Hugging Face, ChromaDB* Apr 2026 – Jun 2026

- Developed a **domain-specific RAG assistant** using transformer models, embeddings, semantic search, and grounded LLM generation.
- Implemented **document ingestion, chunking, retrieval pipelines, context-aware prompting, and similarity search** using ChromaDB.
- Added **retrieval evaluation, response validation, hallucination controls, and regression testing** to measure grounding and reliability.

TECHNICAL SKILLS

Programming: Python, SQL, JavaScript

AI/ML: Machine Learning, NLP, Transformers, PyTorch, scikit-learn, Pandas, NumPy, Model Inference

LLM/RAG: LLMs, Generative AI, RAG, Prompt Engineering, Embeddings, Vector Retrieval, Semantic Search, Similarity Search

Agentic AI: AI Agents, Agentic Workflows, Tool Calling, Task Routing, Workflow Orchestration, Structured Outputs

Frameworks: LangGraph, LangChain, Hugging Face, FastAPI

Vector Databases: ChromaDB, Qdrant, Vector Search, Retrieval Evaluation

Backend/Deployment: REST APIs, Backend Services, Docker, AWS, Git, GitHub Actions, CI/CD

Reliability: Evaluation, Regression Testing, Debugging, Logging, Monitoring, Error Handling, Latency Optimization