

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL
- **AI/ML:** Machine Learning, Deep Learning, Model Inference, Model Evaluation, PyTorch, Hugging Face, Transformers
- **Generative AI:** LLMs, Generative AI, Agentic AI, Multi-Agent Workflows, RAG, LangGraph, LangChain, Prompt Engineering, Context Engineering, Embeddings, Vector Search, Vector Databases
- **Data:** Data Preprocessing, ETL
- **APIs:** FastAPI, REST APIs, HTTP APIs, Async Processing
- **Deployment:** AWS, Docker, CI/CD, ML Deployment, AI Pipeline Deployment
- **Engineering:** Testing, Regression Testing, Evaluation, Error Analysis, Debugging, Performance Optimization, Git, GitHub

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)

June 2026 - Present

 - Built Generative AI workflows using LLM inference, prompt engineering, context engineering, and AI evaluation.
 - Built backend services, REST APIs, and AI pipelines integrating LLM inference, search systems, and automation workflows.
 - Designed RAG pipelines using embeddings and vector search to improve contextual retrieval and AI reliability.
 - Developed AI evaluation pipelines with regression testing, error analysis, validation, and performance monitoring.
 - Applied LangGraph agentic workflows for task orchestration and multi-step LLM applications.
- Data Engineer Intern | flickDONE | [LINK](#)

June 2025 - August 2025

 - Built Python-based ML data pipelines processing 50,000+ records for preprocessing, training, and validation.
 - Engineered automated preprocessing and data validation workflows, reducing manual verification by 70%.
 - Optimized AWS ML pipelines, improving execution speed by 25% through bottleneck analysis.
 - Developed reusable AI workflows for conversational AI, document processing, and business automation.
 - Applied Docker and CI/CD practices to improve deployment reliability, scalability, and maintainability of ML workflows.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)

April 2026 - June 2026

 - Built a Generative AI RAG application using Python, PyTorch, Hugging Face, embeddings, vector search, and LLMs.
 - Designed asynchronous RAG retrieval architecture, improving semantic search efficiency by 50%.
 - Developed modular retrieval pipelines with context management, reducing end-to-end inference latency by 50%.
 - Improved AI output reliability through model evaluation, grounding validation, testing, debugging, and response quality checks.
- AI Search & Recommendation Platform | [LINK](#)

June 2025 - August 2025

 - Designed an AI-powered semantic search platform for retrieval and recommendation workflows.
 - Built FastAPI REST APIs and backend services for ML inference, model serving, semantic search, and orchestration.
 - Integrated Databricks infrastructure, reducing AI routing latency by 35%.
 - Containerized services with Docker, improving deployment reliability and scalability by 25%.
- Morning Market Brief | [LINK](#)

June 2024 - August 2024

 - Built Generative AI workflows using OpenAI APIs, LLM inference, and prompt engineering for financial summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by 45%.
 - Implemented prompt validation, LLM evaluation, and quality checks, improving model reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)

August 2024

 - Built machine learning models using PyTorch and applied ML techniques for model development and evaluation.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with Large Language Models (Coursera) | [CERTIFICATE](#)

March 2024

 - Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows, AI output evaluation, and RAG pipeline performance validation.