

JYOTHIR RAGHAVALU BHOGI

[Jyothir-dev | LinkedIn](#)
[Jyothir-dev \(Jyothir\) \(github.com\)](#)

Email: jyothirraghavalu369@gmail.com
Mobile: +91 6281073800

EDUCATION

Lovely Professional University
Bachelor of Technology (Computer Science & Engineering)

Jalandhar, India
July 2022 - May 2026

SKILLS SUMMARY

- **Programming:** Python, SQL, PySpark
- **Data:** Pandas, NumPy, ETL/ELT
- **AI/ML:** PyTorch, Hugging Face, Transformers, Machine Learning, Deep Learning, NLP, Model Evaluation, Experimentation
- **AI Systems:** LLMs, Agentic AI, Retrieval-Augmented Generation (RAG), Prompt Engineering, LangChain, LlamaIndex, Vector Databases, AI Orchestration, Tool Calling
- **Backend:** FastAPI, REST APIs, Model Inference, Model Deployment, ML Pipelines, CI/CD, Docker
- **Engineering:** Git, GitHub, Debugging, Root Cause Analysis, Performance Optimization

WORK EXPERIENCE

- AI Intern | Know AI | [LINK](#)

June 2026- Present

 - Developed production AI applications using Python, PyTorch, Hugging Face, Transformers, and LLMs.
 - Built backend services and REST APIs supporting LLM inference, scalable AI workflows, and deployment.
 - Designed modular RAG pipelines integrating vector databases to improve retrieval quality and response consistency.
 - Built agent-oriented workflows using prompt engineering and AI orchestration for multi-step LLM tasks.
 - Conducted experimentation, debugging, model evaluation, and performance optimization to improve AI reliability.
- Data Engineer Intern | flickDONE | [LINK](#)

June 2025- August 2025

 - Built scalable NLP pipelines processing 50,000+ datasets for parallel text classification.
 - Engineered automated preprocessing and tokenization workflows, reducing manual verification by **70%**.
 - Optimized AWS-based ML pipelines, improving execution speed by **25%** through bottleneck analysis.
 - Developed reusable workflows supporting conversational AI, intelligent document processing, and agent automation.
 - Collaborated across teams to improve scalability, reliability, and maintainability of production ML systems.

PROJECTS

- Regulatory Compliance RAG Assistant | [LINK](#)

April 26- June 2026

 - Built a RAG system using Python, Hugging Face, ChromaDB, and vector retrieval pipelines.
 - Designed asynchronous retrieval architecture, improving semantic search efficiency by **50%**.
 - Developed modular retrieval and inference pipelines, reducing end-to-end latency by **50%**.
 - Improved zero-shot fraud detection accuracy by **30%** through experimentation and model evaluation.
- AI Search & Recommendation Platform | [LINK](#)

June 25- August 2025

 - Designed an LLM-powered semantic search platform supporting intelligent retrieval and recommendations.
 - Built scalable FastAPI REST APIs for AI inference and orchestration workflows.
 - Integrated Databricks infrastructure, reducing AI routing latency by **35%**.
 - Containerized services with Docker, improving deployment reliability and scalability by **25%**.
- Morning Market Brief | [LINK](#)

June 24- August 2024

 - Built automated AI workflows integrating OpenAI APIs for financial news summarization.
 - Designed embedding generation and vector indexing pipelines, reducing ingestion latency by **45%**.
 - Implemented prompt validation workflows, improving LLM reliability and workflow transparency.
 - Dockerized application components for reproducible deployment and scalable AI automation.

CERTIFICATES

- Python, Data Science & ML (Cipherschools) | [CERTIFICATE](#)

August 2024

 - Built ML models using PyTorch.
 - Developed Dockerized environments supporting reproducible experimentation and deployment.
- Generative AI with LLMs (Coursera) | [CERTIFICATE](#)

March 2024

 - Applied prompt engineering techniques for production LLM applications.
 - Built retrieval workflows and evaluated RAG pipeline performance.