

# LinkedIn Profile Optimization — Jyothir Raghavalu Bhogi

**Target: Generative AI / AI Engineer (entry-level, production-oriented LLM, RAG, agentic AI, ML & data systems)**

## 1. Profile Positioning Strategy

**Primary persona:** Junior Generative AI / AI Engineer who builds production-oriented systems around LLMs — not someone who “experiments with prompts.” Your two internships and three major projects both point the same direction: backend services + orchestration + retrieval + evaluation, wrapped around LLMs. That is a legitimate, narrow, defensible identity — lean into it fully instead of spreading across Data Scientist/Web Dev/Data Analyst framings.

**Secondary target roles (in order of proximity to your evidence):**

RAG Engineer — strongest evidence (Production-Grade RAG Engine, Compliance RAG Assistant, Know AI RAG pipelines)

LLM Engineer — strong (LLM inference, prompt engineering, evaluation across both internships)

AI Agent Engineer — strong (LangGraph orchestrator)

GenAI Application Engineer — strong (all three projects + Know AI)

AI Engineer (generalist) — strong, umbrella term

ML Engineer / MLOps-ML Platform Engineer — moderate (flickDONE pipelines, AWS optimization)

Data Engineer — moderate (flickDONE PySpark/ETL work) — keep as credible secondary, not primary

NLP Engineer — light-moderate (embeddings, semantic search count, but no classic NLP task work like NER/classification — don't overclaim)

**What recruiters should remember about you (one sentence):** *“Junior engineer who has already shipped RAG pipelines, an agent orchestrator, and LLM evaluation workflows behind real APIs — not just notebooks.”*

**Skills to emphasize:** RAG, LLM inference & evaluation, LangChain/LangGraph, embeddings/vector search (Qdrant), FastAPI, Python, AWS, Docker, ML/data pipelines (PySpark), reliability/testing/monitoring framing.

**Skills to de-emphasize (keep listed, don't lead with):** general web development, classic data analysis/BI (SAP Analytics Cloud, Mall Visitor Demographics), older non-AI Python scripting projects, TF-IDF-only recommender work.

**What your current profile is doing wrong:**

No clear headline identity — reads as a generalist B.Tech student rather than a GenAI/LLM specialist.

Older projects (movie recommender, COVID analysis, mall demographics) sit at the same visual priority as your RAG/agent work, diluting the signal.

Internship bullets are likely written as task lists (“worked on X”) rather than achievement-oriented, keyword-rich statements.

No dedicated "Featured" curation — recruiters have to dig for your strongest RAG/agent projects instead of seeing them immediately.

Skills section likely isn't ordered for search relevance (LinkedIn weights the first skills more heavily in search and endorsement prompts).

HOI (Human-Object Interaction) project is computer vision, not GenAI — currently probably positioned too prominently for your target roles.

**What should change:** Headline and About lead with GenAI/LLM/RAG/Agents explicitly. Experience bullets rewritten as outcomes. Projects re-ordered with RAG Engine and LangGraph Orchestrator first. Featured section curated to only the three GenAI projects + GitHub + portfolio. Skills reordered so the first 10 are 100% GenAI/ML-search-relevant. Older CV/analytics projects moved down or archived, not deleted.

## 2. Headline

**Best headline (159 characters):**

Generative AI Engineer | LLM, RAG & AI Agents | Python, LangChain, LangGraph, FastAPI, AWS | Building Production-Oriented ML & Data Systems

**Alternatives:**

AI Engineer (GenAI/LLM/RAG) | Building Agentic AI Systems with LangGraph & LangChain | Python, FastAPI, AWS, Qdrant

Generative AI & ML Engineer | RAG Pipelines, LLM Evaluation, AI Agents | Python | FastAPI | AWS | Docker | Vector Search

Junior AI/ML Engineer specializing in LLM Applications, RAG & Agentic AI | Python · LangChain · FastAPI · AWS

## 3. About Section

Generative AI Engineer building production-oriented systems around large language models — not just experimenting with prompts, but engineering the retrieval, orchestration, and evaluation layers that make LLMs reliable in real applications.

I work across the full GenAI stack: RAG pipelines with embeddings and vector search, multi-step agent orchestration with LangGraph, LLM evaluation and error analysis, and the backend APIs that expose all of it to real users.

At Know AI, I built generative AI workflows combining LLM inference, prompt engineering, and model evaluation, and designed RAG pipelines using embeddings and vector search to improve contextual retrieval and output reliability — backed by FastAPI services for inference, search, and automation. At flickDONE, I built Python-based ML data pipelines processing 50,000+ records, cut manual verification effort by 70% through automated validation workflows, and improved AWS ML pipeline execution speed

by 25% through bottleneck analysis — the data-engineering foundation that production ML/GenAI systems depend on.

My independent projects follow the same pattern — building the system around the model, not just the model:

A stateful multi-agent orchestration platform (LangGraph + LangChain + FastAPI + Gemini) with conditional routing, checkpointing, and failure recovery

A production-style RAG engine (LangChain + Qdrant + AWS) with chunking, embedding, retrieval, reranking, and evaluation

A regulatory-compliance RAG assistant with grounded generation and retrieval quality testing

**Tech stack:** Python · LangChain · LangGraph · FastAPI · PyTorch · Hugging Face · Qdrant (vector search / embeddings) · AWS · Docker · PySpark · SQL · Git

I'm graduating from Lovely Professional University (B.Tech, CSE, 2022–2026) and I'm open to entry-level opportunities in Generative AI Engineering, LLM Engineering, RAG Engineering, AI Agent Engineering, ML Engineering, and MLOps/Data Engineering roles.

## 4. Experience Section

**Know AI Recommended title:** Generative AI Engineer Intern **Description (job description field):** Built production-oriented Generative AI systems — LLM inference workflows, RAG pipelines, and evaluation pipelines — exposed through backend APIs. **Bullets:**

Built Generative AI workflows using LLM inference, NLP, prompt engineering, and structured model evaluation.

Designed RAG pipelines with embeddings and vector search, improving contextual retrieval and output reliability for downstream LLM applications.

Developed backend REST APIs (FastAPI) integrating LLM inference, search, and workflow automation into a single service layer.

Built AI/ML evaluation pipelines covering model testing, error analysis, and output validation to catch regressions before deployment.

Conducted systematic model evaluation and optimization work that directly improved ML reliability across the pipeline.

**flickDONE Recommended title:** Data Engineer Intern **Description:** Built and optimized Python/AWS data pipelines supporting ML training and validation workflows at scale. **Bullets:**

Built Python-based ML data pipelines processing 50,000+ records for preprocessing, training, and validation.

Engineered automated preprocessing and data-validation workflows that reduced manual verification effort by 70%.

Optimized AWS ML pipelines for execution speed, achieving a 25% improvement through targeted bottleneck analysis.

Developed reusable AI workflow components for conversational AI, document processing, and business-process automation.

Collaborated with engineering teams to improve the scalability, reliability, and maintainability of shared ML platforms.

## 5. Featured Section

**Recommended order:**

**Production-Grade RAG Engine** — Title: "Production RAG Engine — LangChain, Qdrant, FastAPI, AWS."

Description: end-to-end ingestion → chunking → embedding → retrieval → reranking → evaluation pipeline built for reliability, not a demo. Recruiter should see: real retrieval evaluation and error handling, not just "I built a chatbot."

**LangGraph Workflow Orchestrator** — Title: "Stateful AI Agent Orchestrator — LangGraph + FastAPI."

Description: multi-step agent workflows with conditional routing, checkpointing, and failure recovery, exposed via REST API. Recruiter should see: agentic AI engineering, not just chained prompts.

**Regulatory Compliance RAG Assistant** — Title: "Compliance RAG Assistant — Grounded Generation over Documents." Description: document ingestion + semantic retrieval + grounded generation with retrieval-quality testing. Recruiter should see: applied RAG in a realistic, constrained domain (compliance), showing judgment about correctness.

**GitHub** — pin your GenAI/RAG/agent repos at the top of your profile README/pinned repos before linking here, so the first thing a click reveals matches your headline.

**Portfolio** (if one exists — otherwise skip; don't fabricate one) — link only if it foregrounds the same three projects.

**Strongest AI-related posts** — only if you have posts already discussing RAG/LLM/agent work; otherwise leave this slot empty rather than featuring an unrelated post.

## 6. Projects Section

**1. Production-Grade RAG Engine** One-line: End-to-end RAG pipeline with retrieval evaluation, built for production reliability rather than a notebook demo. Bullets:

Built document ingestion, chunking, embedding, and vector retrieval pipelines for contextual generation.

Implemented semantic search over a Qdrant vector database.

Added reranking and retrieval evaluation to improve context relevance and response quality.

Implemented fallback execution, monitoring, and error handling to improve pipeline reliability.

Exposed the pipeline through FastAPI for integration into downstream applications. Technologies: Python, FastAPI, LangChain, Qdrant, AWS

**2. LangGraph Workflow Orchestrator** One-line: Stateful multi-agent orchestration platform with conditional routing and failure recovery. Bullets:

Designed multi-step agent workflows with state management and conditional routing.

Built modular components for task decomposition, tool execution, and agent coordination.

Added checkpointing and failure-recovery logic to improve workflow reliability.

Exposed workflow execution through FastAPI REST APIs.

Implemented execution monitoring and testing to validate workflow behavior. Technologies: Python, LangGraph, LangChain, FastAPI, Gemini

**3. Regulatory Compliance RAG Assistant** One-line: Generative AI assistant answering compliance questions with grounded, retrieval-based generation. Bullets:

Designed document ingestion and semantic retrieval pipelines for compliance documents.

Implemented context-grounded generation to reduce unsupported/hallucinated answers.

Built document upload and querying workflows in a Streamlit interface.

Improved response reliability through retrieval testing, evaluation, and debugging. Technologies: Python, Streamlit, LangChain, embeddings, vector search

**Older projects — recommended handling:**

**KEEP (moderate visibility):** Human-Object Interaction Recognition Pipeline — computer vision is adjacent, not central; keep it lower in the list as evidence of broader ML engineering range, not as a headline project.

**MOVE LOWER:** Movie Recommendation System (content-based/collaborative filtering, TF-IDF) — legitimate ML work but not GenAI; useful as ML breadth, not a focus piece.

**MOVE LOWER:** Point Cloud Processing (Open3D) — interesting but unrelated to your target roles; keep as a minor entry.

**REMOVE/ARCHIVE from main profile view (keep on GitHub, don't feature):** COVID-19 data analysis/prediction, Mall Visitor Demographics Analysis, older web-development/social-media conceptual work — these read as coursework/analyst work and actively work against the "GenAI/LLM engineer" positioning if visible near the top.

## 7. Skills Section (Top 30, ordered)

**Tier 1 — Must-have / highest priority (first 10 — drives search visibility):**

Generative AI

Large Language Models (LLM)

Retrieval-Augmented Generation (RAG)

LangChain

LangGraph

AI Agents

Python

Machine Learning

FastAPI

Vector Search

**Tier 2 — Important:** 11. Embeddings 12. Semantic Search 13. Qdrant 14. Prompt Engineering 15. PyTorch 16. Hugging Face 17. Deep Learning 18. Model Evaluation 19. Model Inference 20. AWS

**Tier 3 — Supporting:** 21. Docker 22. Model Deployment 23. PySpark 24. Apache Spark 25. SQL 26. ETL 27. Data Preprocessing 28. REST APIs 29. Git 30. Workflow Automation

## 8. Certifications

### Keep prominently visible:

Generative AI with Large Language Models — Coursera (directly on-brand, keep at the top)

Python, Data Science & Machine Learning — CipherSchools (supports ML/Python positioning)

### Keep but deprioritize (lower in the list):

Python Certification — GUVI & HCL (redundant with the above but not harmful)

### Not particularly useful for target roles (keep on profile but don't surface, don't remove entirely):

Web Development — YHills (works against "not a web developer" positioning; leave it listed but at the very bottom)

SAP Certified Data Analyst — SAP Analytics Cloud (off-thesis for GenAI/ML engineering; keep for breadth but at the bottom, not near the top)

## 9. Education

**Degree title:** Bachelor of Technology, Computer Science and Engineering — Lovely Professional University, Jalandhar, India (2022–2026)

**Description:** CSE undergraduate with a focus on Generative AI, machine learning, and applied software engineering. Independent and internship project work centered on LLM applications, RAG pipelines, agentic AI orchestration, and ML/data pipelines.

**Relevant coursework/skills (only if these were actually part of your coursework — verify before adding):** Data Structures & Algorithms, Machine Learning, Database Management Systems, Cloud Computing — list only what you actually took.

**Activities:** Include only if you have real ones (a GenAI/ML club, hackathon participation, etc.) — don't add generic "coding club" filler if it doesn't exist; an empty activities section is better than a fabricated one.

## 10. Open to Work — Recommended Titles (ranked)

Generative AI Engineer

AI Engineer

LLM Engineer

RAG Engineer

AI Agent Engineer

Machine Learning Engineer

GenAI Application Engineer

NLP Engineer

MLOps Engineer

ML Platform Engineer  
Data Engineer  
Applied AI Engineer  
AI/ML Engineer (entry-level)  
Backend Engineer (AI/ML systems) — reasonable fallback given your FastAPI work  
Junior Software Engineer (Python/AI) — broadest fallback, lowest priority

## 11. Location / Job Preferences

**Location:** Set to a specific city (e.g., your current city in India) rather than leaving it blank — LinkedIn search weights a specified location, and recruiters often filter by proximity even for remote-friendly roles.

**Remote:** Enable — most entry-level GenAI/LLM roles at startups are remote-first in India.

**Hybrid:** Enable.

**On-site:** Enable, and explicitly select major India tech hubs if the "specific locations" field allows multiple entries (Bangalore, Hyderabad, Pune, Gurgaon/NCR, Chennai) — don't restrict yourself to one city given you're an early-career candidate with no location constraints stated.

**India-wide:** Set your Open to Work location scope to allow "India" broadly if LinkedIn's field permits a country-level or multi-city entry, rather than a single city only.

## 12. Profile Photo / Banner

**Profile photo:** Simple, well-lit headshot, neutral or softly blurred background, shoulders-up, direct eye contact, plain professional attire (no logos). Avoid casual selfies or heavily filtered photos — recruiters treat headshot quality as a signal of professionalism.

**Banner:** A clean, dark or deep-navy/charcoal background (reads as "technical," not "corporate stock photo") with a subtle abstract graph/network motif (nodes and edges, faint — evokes embeddings/vector space and agent graphs without being literal or cartoonish).

**Banner text:** "Generative AI · LLMs · RAG · Agentic AI · ML Engineering" — centered or left-aligned, clean sans-serif, one line if it fits, two lines max.

**Visual style:** Minimal, technical, restrained — think "engineering diagram," not "AI hype poster."

**Colors:** Deep navy or charcoal base, one accent color (electric blue, teal, or violet) for the text/motif — avoid neon rainbow AI-cliché gradients.

**What NOT to put on the banner:** No robot/brain/circuit-board clipart, no "AI" in a glowing sci-fi font, no stock photos of people pointing at holograms, no company logos you don't have rights to use, no crowded multi-line taglines.

## 13. 30-Day Content Strategy

Week 1

*Post 1 — RAG:* Topic: "What actually breaks in a RAG pipeline." Hook: "My RAG pipeline had 90% retrieval precision and still gave wrong answers. Here's why." Technical subject: reranking and context relevance vs. raw retrieval accuracy. CTA: "What's the RAG failure mode you've hit most often?" Hashtags: #RAG #LLM #GenerativeAI

*Post 2 — LangGraph:* Topic: why you moved from plain LangChain chains to LangGraph. Hook: "Chains work until your workflow needs to make a decision. Then you need a graph." Technical subject: conditional routing and checkpointing for failure recovery. CTA: "Anyone else find graphs easier to debug than chains?" Hashtags: #LangGraph #AIAgents #GenAI

## Week 2

*Post 3 — Vector databases:* Topic: choosing Qdrant and what changed once you added reranking. Hook: "Adding a reranker fixed more retrieval quality issues than switching embedding models did." Technical subject: embeddings vs. reranking trade-offs. CTA: "Curious what others are using for reranking." Hashtags: #VectorSearch #Qdrant #Embeddings

*Post 4 — LLM evaluation:* Topic: how you built an evaluation pipeline for LLM outputs. Hook: "You can't improve an LLM pipeline you're not measuring." Technical subject: error analysis and output validation approach from Know AI work. CTA: "What metrics do you actually trust for LLM eval?" Hashtags: #LLM #ModelEvaluation #AIEngineering

## Week 3

*Post 5 — FastAPI + LLM serving:* Topic: structuring a FastAPI service around LLM inference and retrieval. Hook: "Most 'LLM app' tutorials skip the API layer. That's the part that actually breaks in production." Technical subject: request handling, timeouts, and fallback logic around model calls. CTA: "What's your API design for LLM-backed endpoints?" Hashtags: #FastAPI #LLM #ProductionAI

*Post 6 — AI agents in practice:* Topic: what "agentic AI" means beyond the buzzword, based on your orchestrator project. Hook: "Agentic AI isn't a bigger prompt. It's state management with extra steps." Technical subject: task decomposition and tool execution in LangGraph. CTA: "How are you defining 'agent' in your own work?" Hashtags: #AIAgents #LangGraph #AgenticAI

## Week 4

*Post 7 — Data pipelines feeding ML:* Topic: the unglamorous data engineering behind ML reliability, from flickDONE. Hook: "Nobody posts about the preprocessing step that cut verification time by 70%. It mattered more than the model." Technical subject: automated validation and preprocessing at scale. CTA: "What's an unglamorous fix that had outsized impact on your pipeline?" Hashtags: #DataEngineering #MLPipelines #PySpark

*Post 8 — Reliability/monitoring:* Topic: what "production-ready" actually means for a GenAI feature. Hook: "A demo and a production RAG system differ mostly in error handling, not model choice." Technical subject: fallback execution, monitoring, and testing across your RAG engine. CTA: "What's your bar for calling something 'production-ready'?" Hashtags: #ProductionAI #AIReliability #RAG

*Post 9 (optional, end of month) — reflection:* Topic: lessons from building three GenAI projects as a student. Hook: "Three GenAI projects later, here's what I'd tell myself at the start." Technical subject: retrieval, orchestration, and evaluation as the recurring hard problems. CTA: "What's the GenAI lesson that took you longest to learn?" Hashtags: #GenerativeAI #AIEngineering #CareerInAI

Avoid opening every post with "I'm excited to share" — vary openers: a direct claim, a specific number, a mini-mistake story, or a technical question.

## 14. LinkedIn Profile SEO Audit

| Keyword                    | Present? | Where to place it                        | Priority |
|----------------------------|----------|------------------------------------------|----------|
| Generative AI Engineer     | Yes      | Headline, About                          | High     |
| GenAI Engineer             | Yes      | Headline (implied), Skills               | High     |
| AI Engineer                | Yes      | Headline, About                          | High     |
| LLM Engineer               | Yes      | Headline, About, Skills                  | High     |
| RAG Engineer               | Yes      | Headline, Projects, Featured             | High     |
| AI Agent Engineer          | Yes      | Headline, Projects                       | High     |
| Machine Learning Engineer  | Yes      | About, Skills, Open to Work              | High     |
| ML Engineer                | Yes      | Skills, Open to Work                     | High     |
| NLP Engineer               | Partial  | Experience (NLP mentioned), Open to Work | Medium   |
| MLOps Engineer             | Partial  | Open to Work, Skills                     | Medium   |
| ML Platform Engineer       | Partial  | Open to Work                             | Medium   |
| GenAI Application Engineer | Yes      | Open to Work, About                      | Medium   |
| Data Engineer              | Yes      | Experience (flickDONE), Open to Work     | Medium   |
| Python                     | Yes      | Headline, About, Skills, Experience      | High     |
| LLM                        | Yes      | Headline, About, Skills                  | High     |
| RAG                        | Yes      | Headline, About, Projects                | High     |
| LangChain                  | Yes      | About, Skills, Projects                  | High     |
| LangGraph                  | Yes      | Headline, Skills, Projects               | High     |
| FastAPI                    | Yes      | Headline, About, Experience, Projects    | High     |
| PyTorch                    | Yes      | Skills                                   | Medium   |
| Hugging Face               | Yes      | Skills                                   | Medium   |

| Keyword          | Present? | Where to place it           | Priority   |
|------------------|----------|-----------------------------|------------|
| AWS              | Yes      | Headline, About, Experience | High       |
| Docker           | Yes      | About, Skills               | Medium     |
| Qdrant           | Yes      | About, Projects, Skills     | High       |
| Vector Database  | Yes      | About (Qdrant), Skills      | High       |
| Embeddings       | Yes      | About, Skills, Projects     | High       |
| Semantic Search  | Yes      | About, Skills               | Medium     |
| Machine Learning | Yes      | Skills, About               | High       |
| Deep Learning    | Yes      | Skills                      | Medium     |
| Generative AI    | Yes      | Headline, About, Skills     | High       |
| Agentic AI       | Yes      | About, Projects, Content    | High       |
| AI Agents        | Yes      | Headline, Skills, Projects  | High       |
| ML Pipelines     | Yes      | About, Experience           | Medium     |
| Data Pipelines   | Yes      | Experience, Skills          | Medium     |
| Model Evaluation | Yes      | Experience, Skills          | Medium     |
| Model Inference  | Yes      | Experience, Skills          | Medium     |
| Model Deployment | Yes      | Skills                      | Low-Medium |
| PySpark          | Yes      | Experience, Skills          | Medium     |
| Apache Spark     | Yes      | Skills                      | Low-Medium |
| SQL              | Yes      | Skills                      | Medium     |

## 15. Final Profile (Copy-Paste Ready)

**NAME** Jyothir Raghavalu Bhogi

**HEADLINE** Generative AI Engineer | LLM, RAG & AI Agents | Python, LangChain, LangGraph, FastAPI, AWS | Building Production-Oriented ML & Data Systems

**ABOUT** Generative AI Engineer building production-oriented systems around large language models — not just experimenting with prompts, but engineering the retrieval, orchestration, and evaluation layers that make LLMs reliable in real applications.

I work across the full GenAI stack: RAG pipelines with embeddings and vector search, multi-step agent orchestration with LangGraph, LLM evaluation and error analysis, and the backend APIs that expose all of it to real users.

At Know AI, I built generative AI workflows combining LLM inference, prompt engineering, and model evaluation, and designed RAG pipelines using embeddings and vector search to improve contextual retrieval and output reliability — backed by FastAPI services for inference, search, and automation. At flickDONE, I built Python-based ML data pipelines processing 50,000+ records, cut manual verification effort by 70% through automated validation workflows, and improved AWS ML pipeline execution speed by 25% through bottleneck analysis — the data-engineering foundation that production ML/GenAI systems depend on.

My independent projects follow the same pattern — building the system around the model, not just the model:

A stateful multi-agent orchestration platform (LangGraph + LangChain + FastAPI + Gemini) with conditional routing, checkpointing, and failure recovery

A production-style RAG engine (LangChain + Qdrant + AWS) with chunking, embedding, retrieval, reranking, and evaluation

A regulatory-compliance RAG assistant with grounded generation and retrieval quality testing

Tech stack: Python · LangChain · LangGraph · FastAPI · PyTorch · Hugging Face · Qdrant (vector search / embeddings) · AWS · Docker · PySpark · SQL · Git

I'm graduating from Lovely Professional University (B.Tech, CSE, 2022–2026) and I'm open to entry-level opportunities in Generative AI Engineering, LLM Engineering, RAG Engineering, AI Agent Engineering, ML Engineering, and MLOps/Data Engineering roles.

## EXPERIENCE

*Know AI — Generative AI Engineer Intern (June 2026 – August 2026)* Built production-oriented Generative AI systems — LLM inference workflows, RAG pipelines, and evaluation pipelines — exposed through backend APIs.

Built Generative AI workflows using LLM inference, NLP, prompt engineering, and structured model evaluation.

Designed RAG pipelines with embeddings and vector search, improving contextual retrieval and output reliability for downstream LLM applications.

Developed backend REST APIs (FastAPI) integrating LLM inference, search, and workflow automation into a single service layer.

Built AI/ML evaluation pipelines covering model testing, error analysis, and output validation to catch regressions before deployment.

Conducted systematic model evaluation and optimization work that directly improved ML reliability across the pipeline.

*flickDONE — Data Engineer Intern (June 2025 – August 2025)* Built and optimized Python/AWS data pipelines supporting ML training and validation workflows at scale.

Built Python-based ML data pipelines processing 50,000+ records for preprocessing, training, and validation.

Engineered automated preprocessing and data-validation workflows that reduced manual verification effort by 70%.

Optimized AWS ML pipelines for execution speed, achieving a 25% improvement through targeted bottleneck analysis.

Developed reusable AI workflow components for conversational AI, document processing, and business-process automation.

Collaborated with engineering teams to improve the scalability, reliability, and maintainability of shared ML platforms.

**EDUCATION** Lovely Professional University — Bachelor of Technology, Computer Science and Engineering (2022–2026), Jalandhar, India CSE undergraduate with a focus on Generative AI, machine learning, and applied software engineering. Independent and internship project work centered on LLM applications, RAG pipelines, agentic AI orchestration, and ML/data pipelines.

## PROJECTS

*Production-Grade RAG Engine* End-to-end RAG pipeline with retrieval evaluation, built for production reliability rather than a notebook demo.

Built document ingestion, chunking, embedding, and vector retrieval pipelines for contextual generation. Implemented semantic search over a Qdrant vector database.

Added reranking and retrieval evaluation to improve context relevance and response quality.

Implemented fallback execution, monitoring, and error handling to improve pipeline reliability.

Exposed the pipeline through FastAPI for integration into downstream applications. Technologies: Python, FastAPI, LangChain, Qdrant, AWS

*LangGraph Workflow Orchestrator* Stateful multi-agent orchestration platform with conditional routing and failure recovery.

Designed multi-step agent workflows with state management and conditional routing.

Built modular components for task decomposition, tool execution, and agent coordination.

Added checkpointing and failure-recovery logic to improve workflow reliability.

Exposed workflow execution through FastAPI REST APIs.

Implemented execution monitoring and testing to validate workflow behavior. Technologies: Python, LangGraph, LangChain, FastAPI, Gemini

*Regulatory Compliance RAG Assistant* Generative AI assistant answering compliance questions with grounded, retrieval-based generation.

Designed document ingestion and semantic retrieval pipelines for compliance documents.

Implemented context-grounded generation to reduce unsupported/hallucinated answers.

Built document upload and querying workflows in a Streamlit interface.

Improved response reliability through retrieval testing, evaluation, and debugging. Technologies: Python, Streamlit, LangChain, embeddings, vector search

*(Lower priority, keep visible below the above three): Human-Object Interaction Recognition Pipeline (Python, OpenCV, YOLOv8, MediaPipe, FastAPI); Movie Recommendation System (TF-IDF, content-based/collaborative filtering, Streamlit); Point Cloud Processing (Open3D)*

## **CERTIFICATIONS** (in this order)

Generative AI with Large Language Models — Coursera  
Python, Data Science & Machine Learning — CipherSchools  
Python Certification — GUVI & HCL  
Web Development — YHills  
SAP Certified Data Analyst — SAP Analytics Cloud

**SKILLS** (top 10 shown, full 30 above in Section 7) Generative AI, Large Language Models (LLM), Retrieval-Augmented Generation (RAG), LangChain, LangGraph, AI Agents, Python, Machine Learning, FastAPI, Vector Search

## **FEATURED**

Production-Grade RAG Engine  
LangGraph Workflow Orchestrator  
Regulatory Compliance RAG Assistant  
GitHub profile (pin GenAI/RAG/agent repos at top first)

**OPEN TO WORK** Generative AI Engineer, AI Engineer, LLM Engineer, RAG Engineer, AI Agent Engineer, Machine Learning Engineer, GenAI Application Engineer, NLP Engineer, MLOps Engineer, ML Platform Engineer, Data Engineer — Location: India-wide (Remote / Hybrid / On-site all enabled)

**BANNER** Deep navy/charcoal background, faint abstract node-and-edge graph motif, single-accent color text: "Generative AI · LLMs · RAG · Agentic AI · ML Engineering"